


Workshop: AgentOps

Evaluate, Ship, Observe, and Operate Microsoft Foundry agents



Workshop

Duration: 6h hands-on; 4h demo
Structure: Three modules
Difficulty: 300: Advanced
Format: Remote or onsite

Description

Learn how to apply AgentOps to your Microsoft Foundry projects, from evaluating an agent to releasing, monitoring, and improving it. Part of the AI Governance Value Based Delivery, this workshop combines practical examples, discussion, and guided activities.

Audience

Developers, architects, platform engineers, and governance teams working with AI agents. Familiarity with Azure and agent development is recommended.

Outcomes

- Apply AgentOps practices to your own project.
- Choose evaluation metrics and acceptance criteria.
- Configure tracing to capture agent behavior.
- Evaluate agent responses and actions.
- Automate evaluation and deployment.
- Monitor agent quality and performance.
- Set up continuous evaluation.
- Respond to failures and improve the agent.

Format

- **Hands-on:** 60m presentation and 60m lab per module.
- **Demo:** 60m presentation and 20m instructor demonstration per module.

Take all three modules or select those relevant to your team. Modules and format are agreed before the workshop.

Prerequisites

- **Hands-on:** a computer with the tools and access described in [participant pre-work](#).
- **Demo:** no participant Azure access is needed.

The learning environment and participant setup are prepared before the session.

Scope

- **Evaluate:** Foundry evaluators; test datasets; rubric evaluation; red teaming.
 - **Ship:** Versioning; CI/CD; release gates and approvals; rollback.
 - **Observe and Operate:** Tracing; monitoring; continuous evaluation; alerts and notifications; incident response.
- Observe and Operate are distinct practices taught together.

Methodology

- **Learn:** practical presentations and discussion.
- **Apply:** guided labs or demos in Microsoft Foundry.
- **Reinforce:** review results and connect the four practices.

Agenda

Module	Duration	Learning artifact
Evaluate	2h hands-on 1h 20m demo	Sample evaluation criteria
Ship	2h hands-on 1h 20m demo	Sample release checklist
Observe and Operate	2h hands-on 1h 20m demo	Sample runtime signal review

Workshop materials

- Presentation materials for the selected modules.
- Guided lab instructions and sample scenarios.
- Examples of evaluation criteria, release checks, and operational findings.

The workshop uses sample scenarios, not your production environment. Separate implementation guidance helps your team apply these practices to its own projects.

Optional advanced module

Additional 2h hands-on (level 400).

Responsible AI in practice: ASSERT; Agent Control Specification (ACS); Agent hooks.