


Workshop: AgentOps

Evaluate, Ship, Observe, and Operate Microsoft Foundry agents



Workshop

Duration: 6h hands-on; 4h demo
Structure: Select one mode before delivery
Difficulty: 300: Advanced
Format: Remote or onsite

Description

Apply Evaluate, Ship, Observe, and Operate to Microsoft Foundry agents through one preselected mode: participant hands-on labs or instructor-led demonstrations. Intended for teams that govern, build, release, and operate AI agents.

Outcomes

- Explain the four AgentOps practices and how they connect.
- Recognize evaluation evidence and release readiness controls.
- Interpret runtime signals and operational responses.
- Relate the practices to guided Microsoft Foundry scenarios.

Prerequisites

- **Hands-on:** complete [participant pre-work](#).
- **Instructor/admin:** accept the pre-work readiness gate and prepare evidence. Demo attendees need no Azure access. Setup is outside workshop time.

Evaluate facilitation

After the 60-minute presentation, select the 60-minute lab or 20-minute demo.

Hands-on allocation

Minutes	Participant activity and evidence
0-5	Tool roles, handoff and workspace context
5-13	Tools, eight cases, mappings and thresholds
13-25	One CLI run, progress and baseline report
25-35	Rows, completeness and tool/source failures
35-44	Domain rubric calibration and native safety
44-51	Conversation scope and red-team coverage
51-60	Decision and evidence handoff to Ship

Selected 20-minute demo

Allow **3m** for azd/Accelerator roles, target, dataset and config; **3m** to start the CLI run and open a matching rehearsal bundle; **5m** for report, raw rows, baseline and both defects; **5m** for native rubric, safety, conversation and red-team evidence; **4m** for the decision, gaps and handoff.

Cloud completion is not guaranteed. Label pre-completed results **instructor evidence review**, never the in-progress run. Follow lab failure handling; do not resubmit blindly or present missing evidence as a pass.

Methodology

- **Learn:** presentations, demos, and discussion.
- **Apply:** guided labs or instructor walkthroughs in Microsoft Foundry.
- **Reinforce:** review results and connect the four practices.

Scope

- **Evaluate:** define criteria and interpret evaluation evidence.
- **Ship:** establish versioning, release gates, and readiness.
- **Observe:** interpret runtime traces, metrics, logs, and usage.
- **Operate:** respond, maintain, manage change, and improve. Observe and Operate remain distinct practices in one shared module.

Agenda

Module	Duration	Learning artifact
Evaluate	2h hands-on 1h 20m demo	Sample evaluation criteria
Ship	2h hands-on 1h 20m demo	Sample release checklist
Observe and Operate	2h hands-on 1h 20m demo	Sample runtime signal review

Delivery options

- **Hands-on (preferred):** 60m presentation and 60m participant lab per module.
- **Demo (fallback):** 60m presentation and 20m instructor demo per module.
- **Advanced extension (optional):** two hours hands-on or one hour demo at difficulty 400, producing a tested runbook, regression case, and updated gate.

Choose the mode before delivery. Do not provision environments during the workshop.

Preparation and delivery

1. Confirm participants, selected modules, and outcomes.
2. Rehearse the selected mode; use labelled evidence discussion if readiness is incomplete.
3. Capture questions, takeaways, and follow-up learning resources.

Reference implementation

Evaluate uses the released [AgentOps Accelerator](#) public CLI. Native Microsoft Foundry activities supplement its gate without requiring internal Accelerator APIs.